



Research

Our research teams investigate the safety, inner workings, and societal impacts of AI models – so that artificial intelligence has a positive impact as it becomes increasingly capable.

Research teams: [Alignment](#) [Economic Research](#) [Interpretability](#) [Societal Impacts](#)

Interpretability

The mission of the Interpretability team is to discover and understand how large language models work internally, as a foundation for AI safety and positive outcomes.

Alignment

The Alignment team works to understand the risks of AI models and develop ways to ensure that future ones remain helpful, honest, and harmless.

Societal Impacts

Working closely with the Anthropic Policy and Safeguards teams, Societal Impacts is a technical research team that explores how AI is used in the real world.

Frontier Red Team

The Frontier Red Team analyzes the implications of frontier AI models for cybersecurity, biosecurity, and autonomous systems.

1	2	9	0	3	8	4	7	5	8	2	4	5	1	7	3	0	5	1	2	9
6	8	4	9	5	7	1	2	4	1	9	2	6	3	7	9	4	3	6	8	4
9	5	8	1	7	6	3	4	1	8	2	5	3	6	8	7	5	2	9	5	8
1	7	3	2	6	9	5	8	9	4	1	6	2	3	8	5	7	4	1	7	3
7	3	2	8	4	5	1	9	3	8	4	1	2	6	9	4	5	6	7	3	2
8	1	6	3	5	4	2	7	6	7	2	8	1	9	3	5	8	9	8	1	6
5	4	7	4	9	1	8	6	4	2	7	6	9	3	1	4	5	2	5	4	7
7	3	9	2	1	8	6	5	3	9	4	8	2	7	1	6	3	4	7	3	9
2	7	1	5	3	0	4	6	0	5	8	9	4	7	1	6	8	2	7	1	
2	5	1	8	9	3	4	0	0	3	8	5	7	2	9	4	0	7	2	5	1
8	1	5	6	7	2	8	4	0	1	9	2	8	3	7	4	6	3	9	1	5
9	1	5	6	7	2	8	4	0	1	9	2	8	3	7	4	6	3	9	1	5

What if we could read
an AI model's thoughts?

What if we could read an AI model's thoughts?

Natural Language Autoencoders: Turning Claude's thoughts into text

Interpretability May 7, 2026

AI models like Claude talk in words but think in numbers. In this study, we train Claude to translate its thoughts into human-readable text.

Alignment May 8, 2026

Teaching Claude why

New research on how we've reduced agentic misalignment.

Research Apr 24, 2026

Project Deal

We created a marketplace for employees in our San Francisco office, with one big twist. We tasked Claude with buying, selling and negotiating on our colleagues' behalf.

Societal Impacts Mar 18, 2026

What 81,000 people want from AI

We invited Claude.ai users to share how they use AI, what they dream it could make possible, and what they fear it might do. Nearly 81,000 people participated—the largest and most multilingual qualitative study of its kind. Here's what we found.

Policy Dec 18, 2025

Project Vend: Phase two

In June, we revealed that we'd set up a small shop in our San Francisco office lunchroom, run by an AI shopkeeper. It was part of Project Vend, a free-form experiment exploring how well AIs could do on complex, real-world tasks. How has Claude's business been since we last wrote?

Publications

Economic Research May 27, 2026

Coding agents in the social sciences

Announcements May 22, 2026

Project Glasswing: An initial update

Policy May 14, 2026

2028: Two scenarios for global AI leadership

Alignment May 8, 2026

Teaching Claude why

Interpretability May 7, 2026

Natural Language Autoencoders: Turning Claude's thoughts into text

Alignment May 7, 2026

Donating our open-source alignment tool

Policy May 7, 2026

Focus areas for The Anthropic Institute

Societal Impacts Apr 30, 2026

How people ask Claude for personal guidance

Science Apr 29, 2026

Evaluating Claude’s bioinformatics research capabilities with BioMysteryBench

Economic Research Apr 22, 2026

Announcing the Anthropic Economic Index Survey

See more ↓

Join the Research team

See open roles →



Products

- Claude
- Claude Code
- Claude Code Enterprise
- Claude Cowork
- Claude Security
- Claude for Chrome
- Claude for Slack
- Claude for Microsoft 365
- Skills
- Max plan

Team plan

Enterprise plan

Download app

Pricing

Log in to Claude

Models

Mythos Preview

Opus

Sonnet

Haiku

Solutions

AI agents

Code modernization

Coding

Customer support

Education

Financial services

Government

Healthcare

Legal

Life sciences

Nonprofits

Security

Small business

Claude Platform

Overview

Developer docs

Pricing

[Marketplace](#)

[Regional compliance](#)

[Claude on AWS](#)

[Google Cloud's Vertex AI](#)

[Microsoft Foundry](#)

[Console login](#)

Resources

[Blog](#)

[Claude partner network](#)

[Community](#)

[Connectors](#)

[Courses](#)

[Customer stories](#)

[Engineering at Anthropic](#)

[Events](#)

[Inside Claude Code](#)

[Inside Claude Cowork](#)

[Inside Claude Enterprise](#)

[Inside Claude Security](#)

[Plugins](#)

[Powered by Claude](#)

[Service partners](#)

[Startups program](#)

[Tutorials](#)

[Use cases](#)

Help and security

[Availability](#)

[Status](#)

[Support center](#)

Company

[Anthropic](#)

[Careers](#)

[Economic Futures](#)

[Research](#)

[News](#)

[Claude's Constitution](#)

[Responsible Scaling Policy](#)

[Security and compliance](#)

[Transparency](#)

Terms and policies

[Privacy choices](#)

[Privacy policy](#)

[Consumer health data privacy policy](#)

[Responsible disclosure policy](#)

[Terms of service: Commercial](#)

[Terms of service: Consumer](#)

[Usage policy](#)

© 2026 Anthropic PBC

